



CS 366

Intro to Computer Security

Dr. Stella Sun
EECS
University of Tennessee
Fall 2022

Today's Class

- Buffer overflow defenses

What Do Attacks Have in Common

- ❖ Attacker is able to input some data used by the program
- ❖ This causes unintentional access to some memory area in the program, e.g.,
 - past a buffer
 - arbitrary locations on the stack

Defense Category 1: Making Exploitation Harder

- ❖ Recall the steps of a stack smashing
 - injecting attacker code into memory (no zeros)
 - overwrite %eip to point to and run attacker code
 - finding the return address (good guess)

How Can We Make These Steps ~~Harder~~

- ❖ Basic idea
 - complicating exploitation by changing compiler, libraries and/or OS (i.e., architectural design)
 - so that the application code doesn't need to be changed
- ❖ Let's look at some techniques
 - stack canaries
 - ASLR

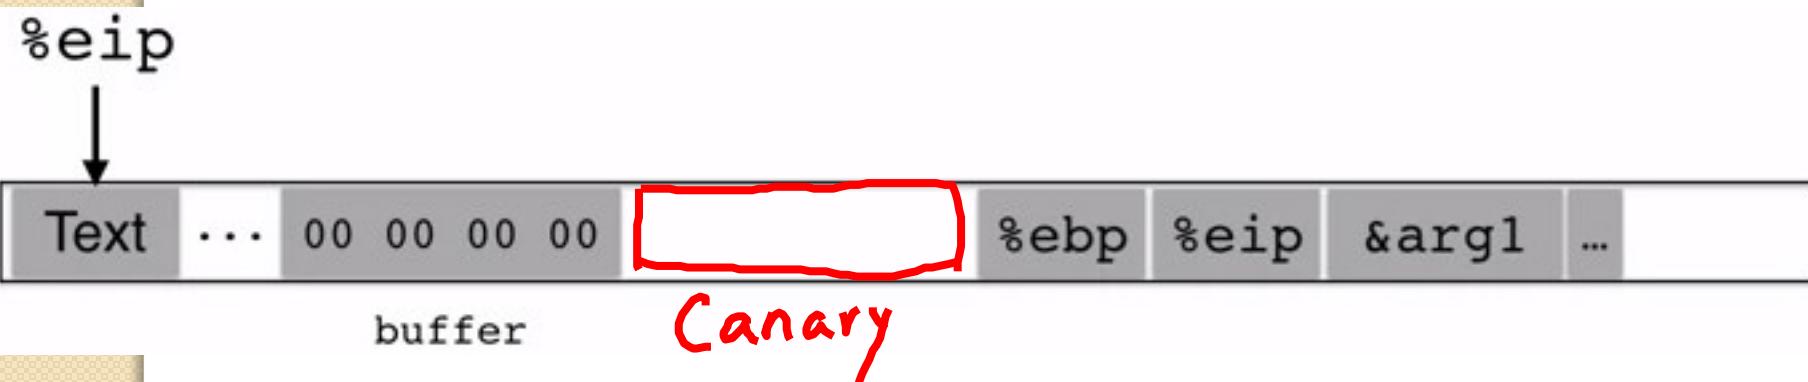
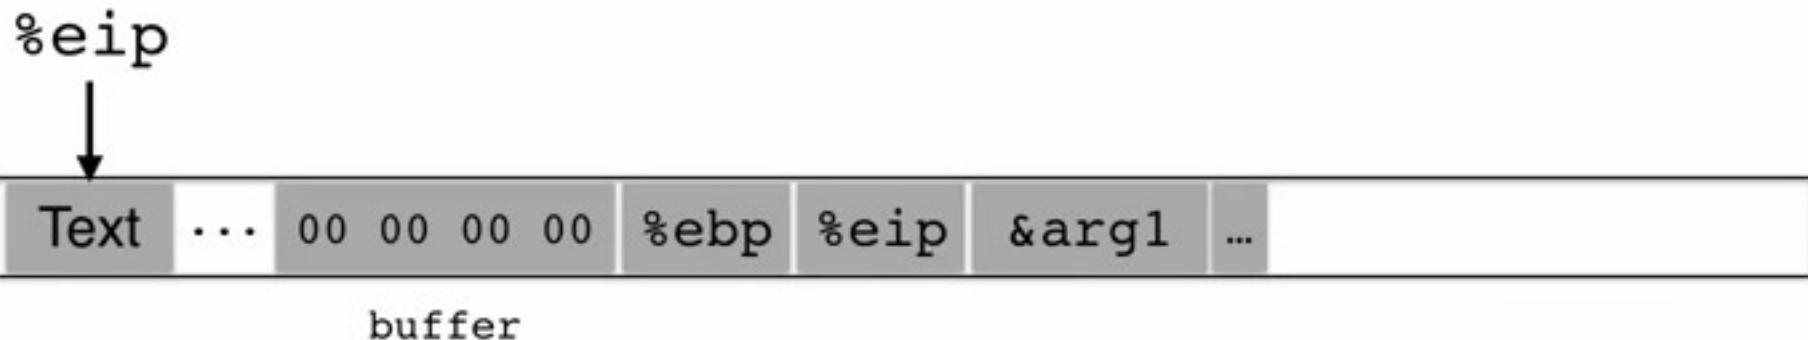
Stack Canaries

- ❖ Why called canary?



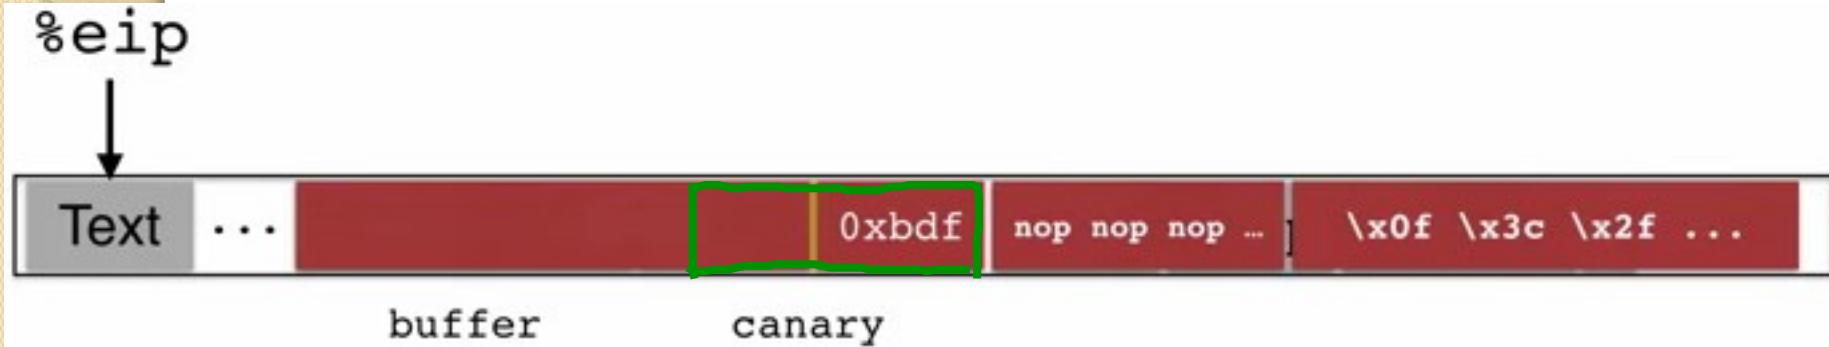
Stack Canaries

- ❖ Detecting overflow with canaries



Stack Canaries

- ❖ If the attacker overruns the buffer



- ❖ The attacker will also need to overrun the canary to overwrite %eip
- ❖ We can change the compiler such that before returning, it checks if the canary value is changed

What Value Should The Canary Have

- ❖ Basic idea
 - choose one that's harder to guess or use
- ❖ Terminator canaries
 - NULL, CR, LF, etc.
 - leveraging the fact that string functions stop copying there
 - attackers cannot use string functions as attack vector

What Value Should The Canary Have

- ❖ Random canaries
 - write a new random value at each process start
 - store this value in a global variable and write-protect it
 - insert into each stack frame
 - known only to the buffer overflow protection code
 - attackers need to be able to steal (read) the canary value

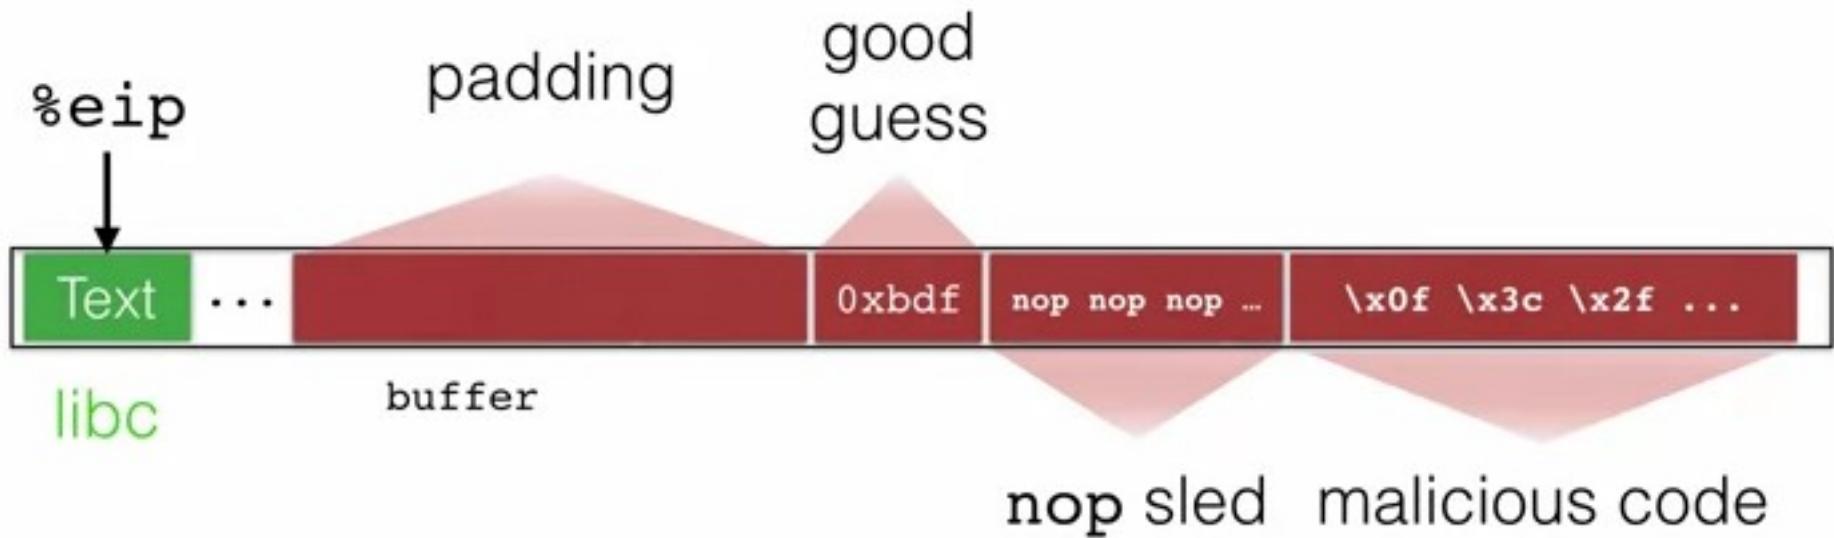
What Value Should The Canary Have

- ❖ Random XOR canaries
 - random canaries XORed with control data (e.g., %eip)
 - slightly better because in some cases the %eip can be changed without overrunning the canary (e.g., when a pointer is overflowed to point to a %eip)
 - attackers needs to be able to steal (read) the canary value
 - only protect control data (or whatever is being XORed), not other data or pointer itself

What If Canaries Fail

- ❖ For example: attacker successfully modified %eip without changing the canary
 - then attacker can still get his injected code to run
- ❖ Solution: make the stack non-executable
 - general idea: make everything other than the code (or text) region non-executable
 - called data execution protection (DEP) or W^X
 - very widely used memory protection

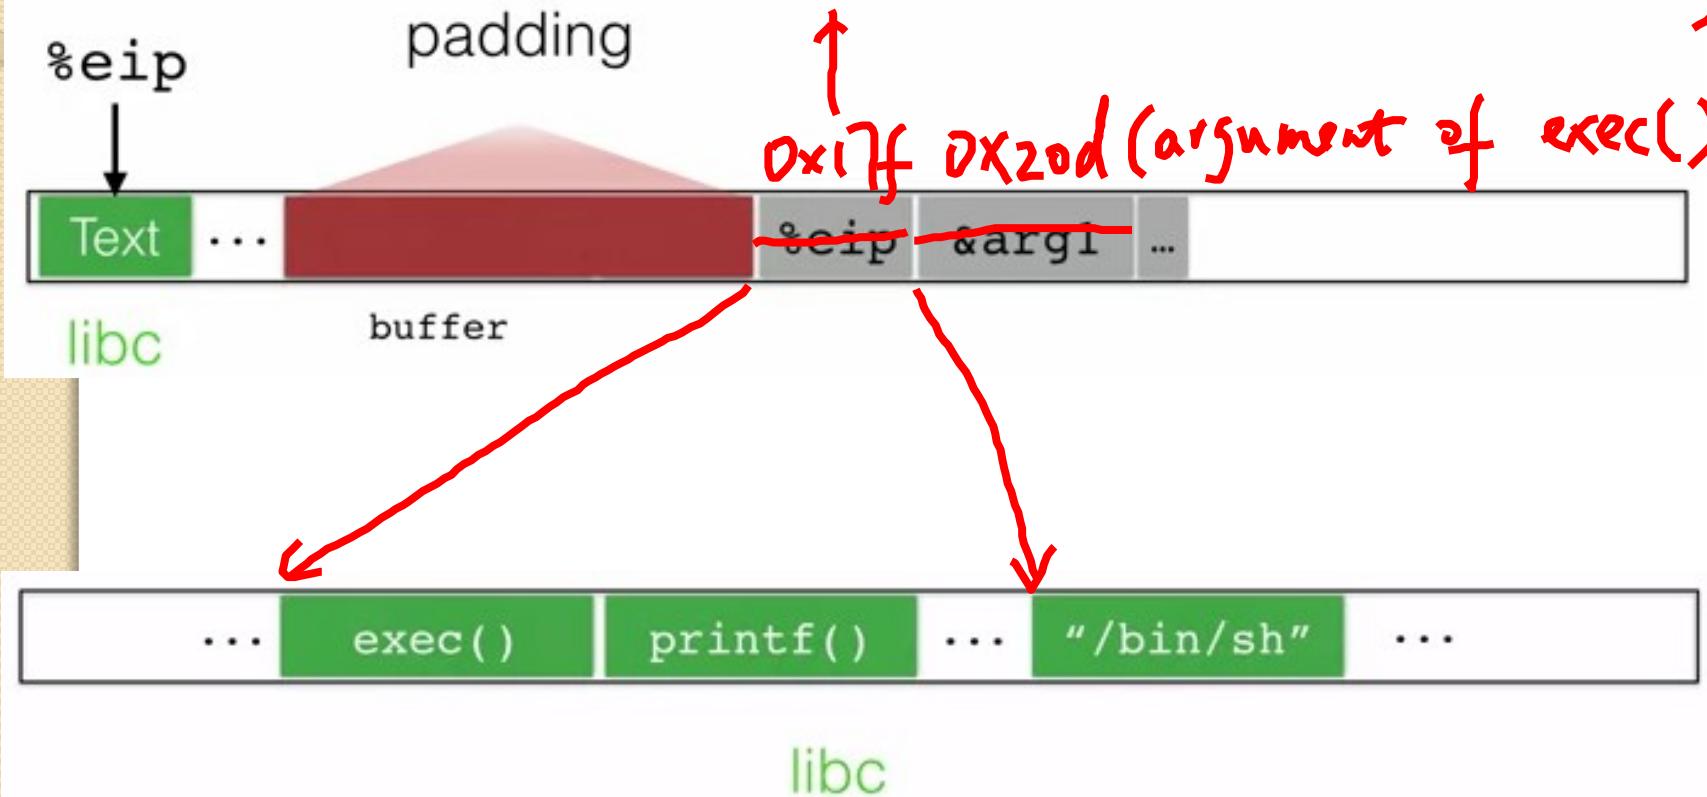
Unfortunately, Can't Prevent return ~~libc~~



Unfortunately, Can't Prevent return libc

- ❖ Instead...

(Known location with code attacker wants)

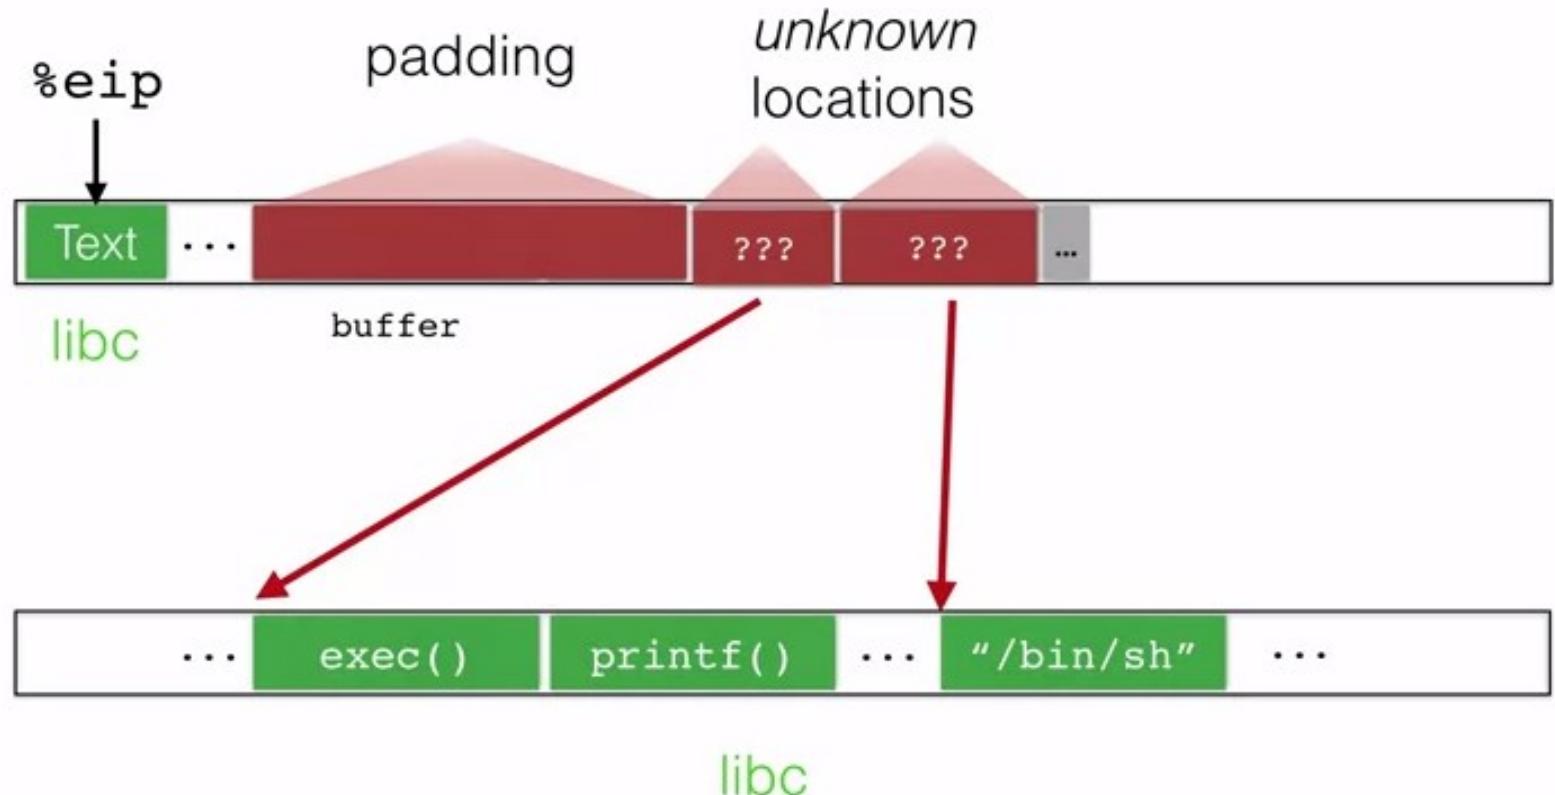


- ❖ Upon returning from the overrun function, the current process is turned into a shell

Use Address Space Layout Randomization (ASLR)

- ❖ Basic idea:

- randomly put standard libraries and other elements (e.g., stack) in memory to make it harder to guess



ASLR

- ❖ Available on most modern OSs
- ❖ Downside
 - only shifts the base of memory areas (not the locations within those areas)
 - may not apply to program code, only libraries
 - needs sufficient randomness, or can be brute forced

took 216 seconds to brute force ASLR on 32-bit systems
[Shacham et al., "On the effectiveness of address-space randomization," (CCS'04)]

- fortunately, 64-bits systems are hard to brute force

Return-Oriented Programming ~~(ROP)~~

- ❖ A special class of return-to-libc attack that is gaining popularity

[Schacham et al., The Geometry of Innocent Flesh on the Bone: Return-to-libc without Function Calls (on the x86), CCS'07]

- ❖ Basic idea:
 - Instead of the entire function, it pieces together code that's not intended to be there that ends with ret (called gadgets)
 - Harder to detect and avoid
 - Intuition: one can make out more words than were intentionally put there, e.g., “dress” is a suffix of “address”, “head” can be made out from “the address”
 - This is possible with dense languages such as x86

Dense x86 Instructions

- ❖ Suppose we have two instructions in libc:

f7 c7 07 00 00 00	test \$0x00000007, %edi
0f 95 45 c3	setnz -61(%ebp)

- ❖ Starting one byte later, the attacker instead obtains:

c7 07 00 00 00 0f	movl \$0x0f000000, (%edi)
95	xchg %ebp, %eax
45	inc %ebp
c3	ret

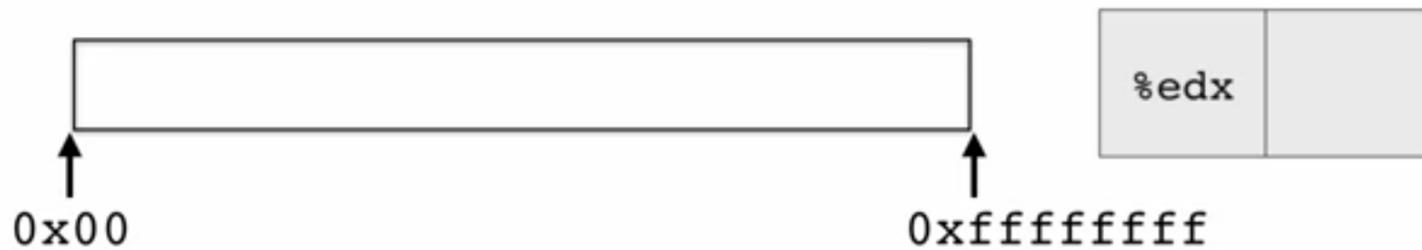
- ❖ How frequently does this happen?

- Very frequently with dense languages
- The set of gadgets are Turing complete: ROP can do anything possible with x86 code

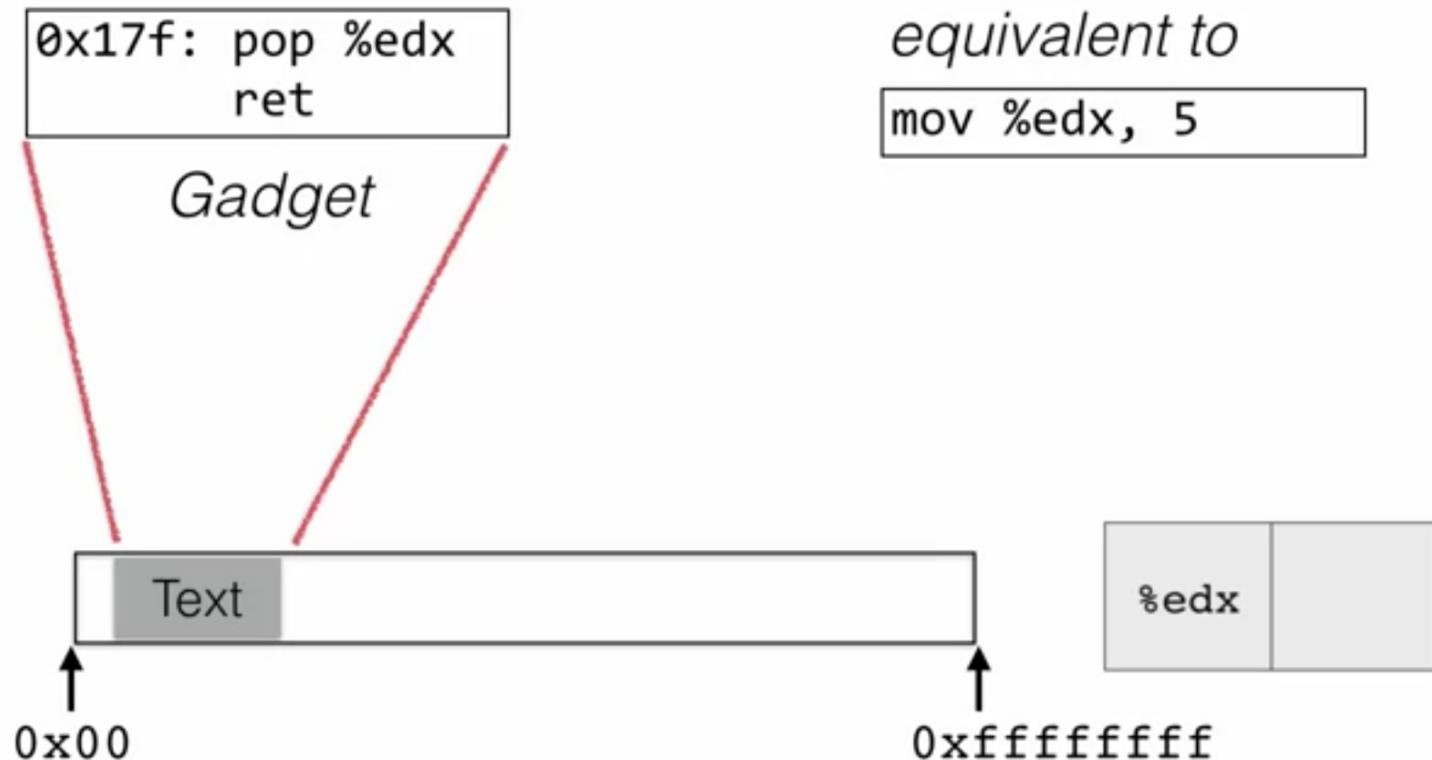
Simple ROP Example

equivalent to

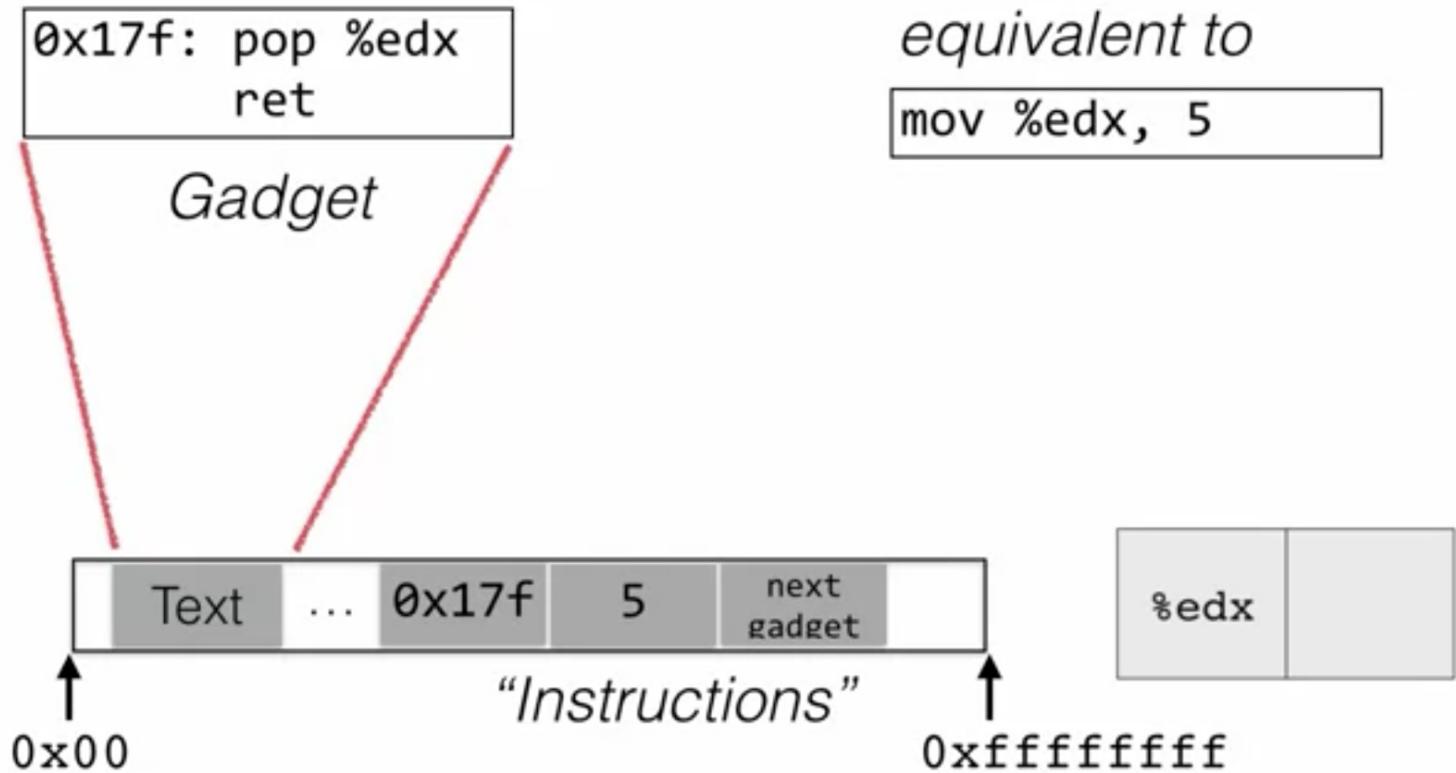
```
mov %edx, 5
```



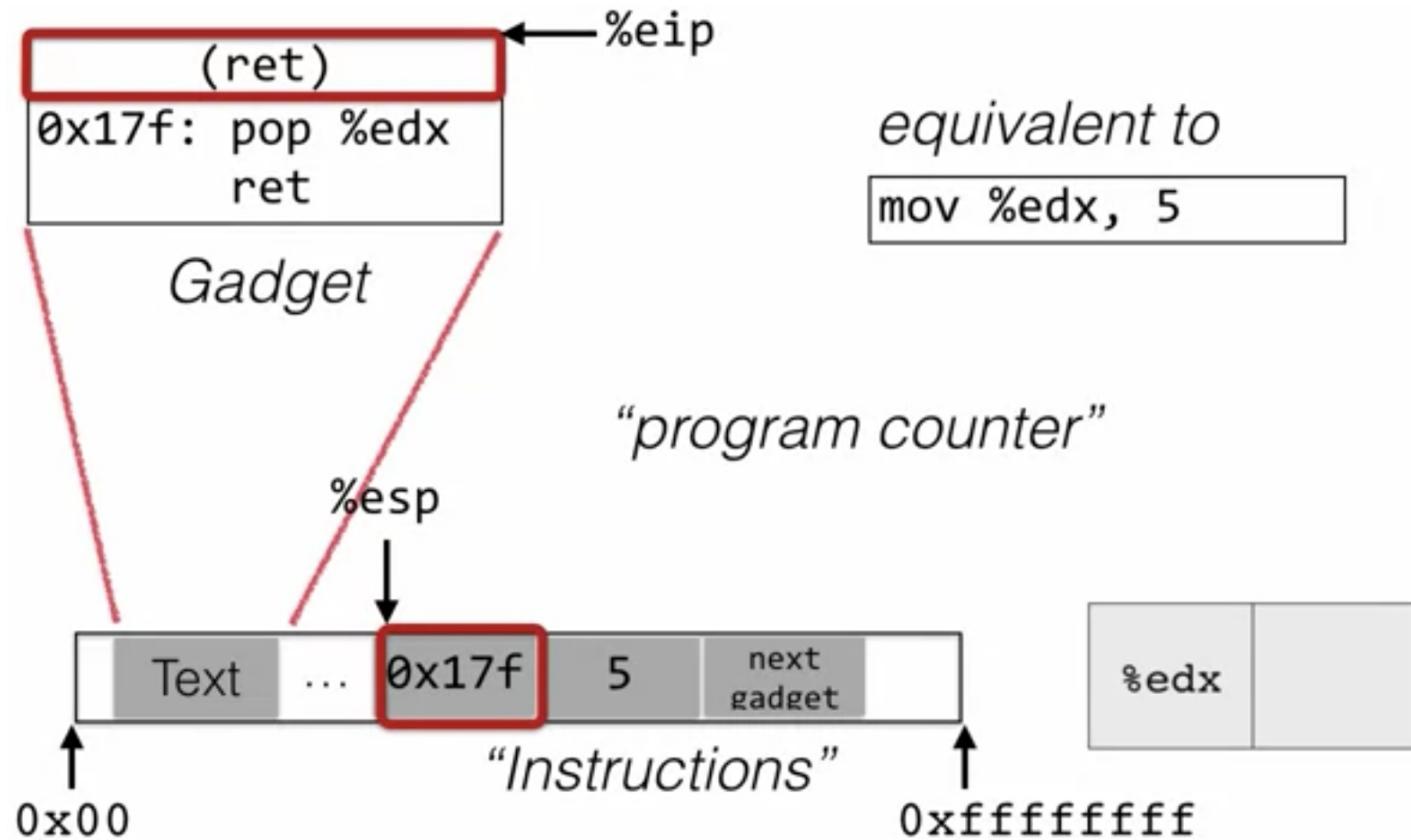
Simple ROP Example



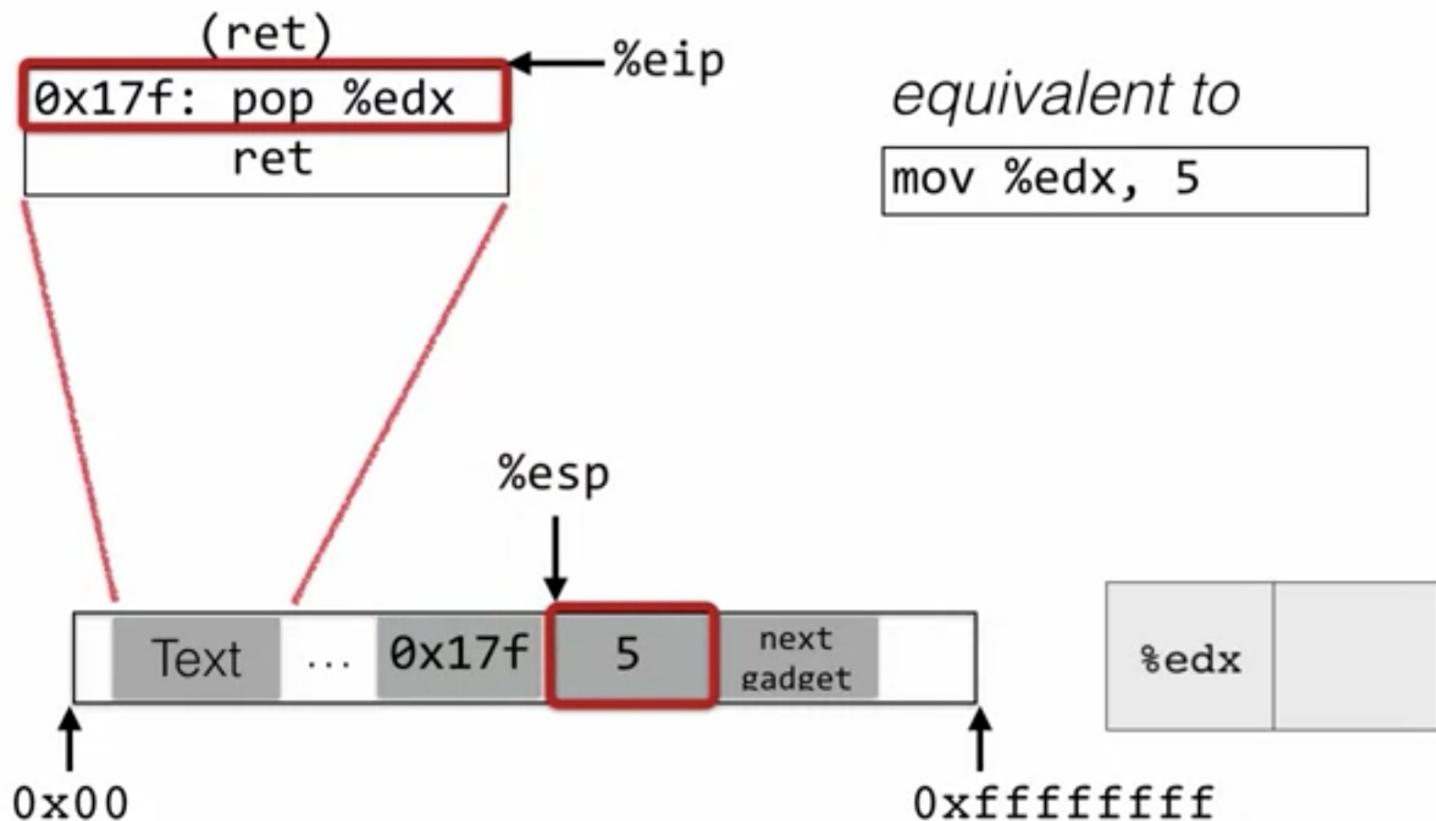
Simple ROP Example



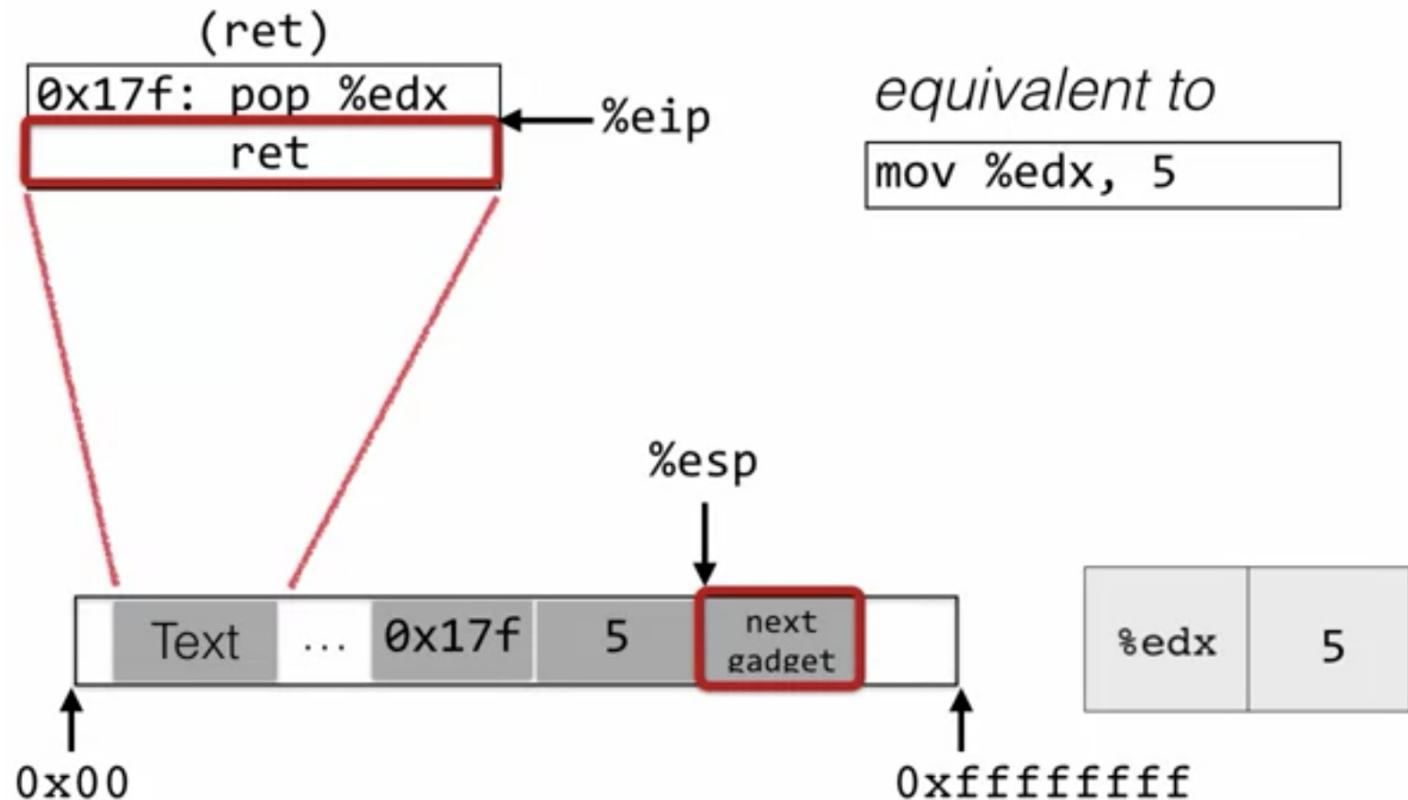
Simple ROP Example



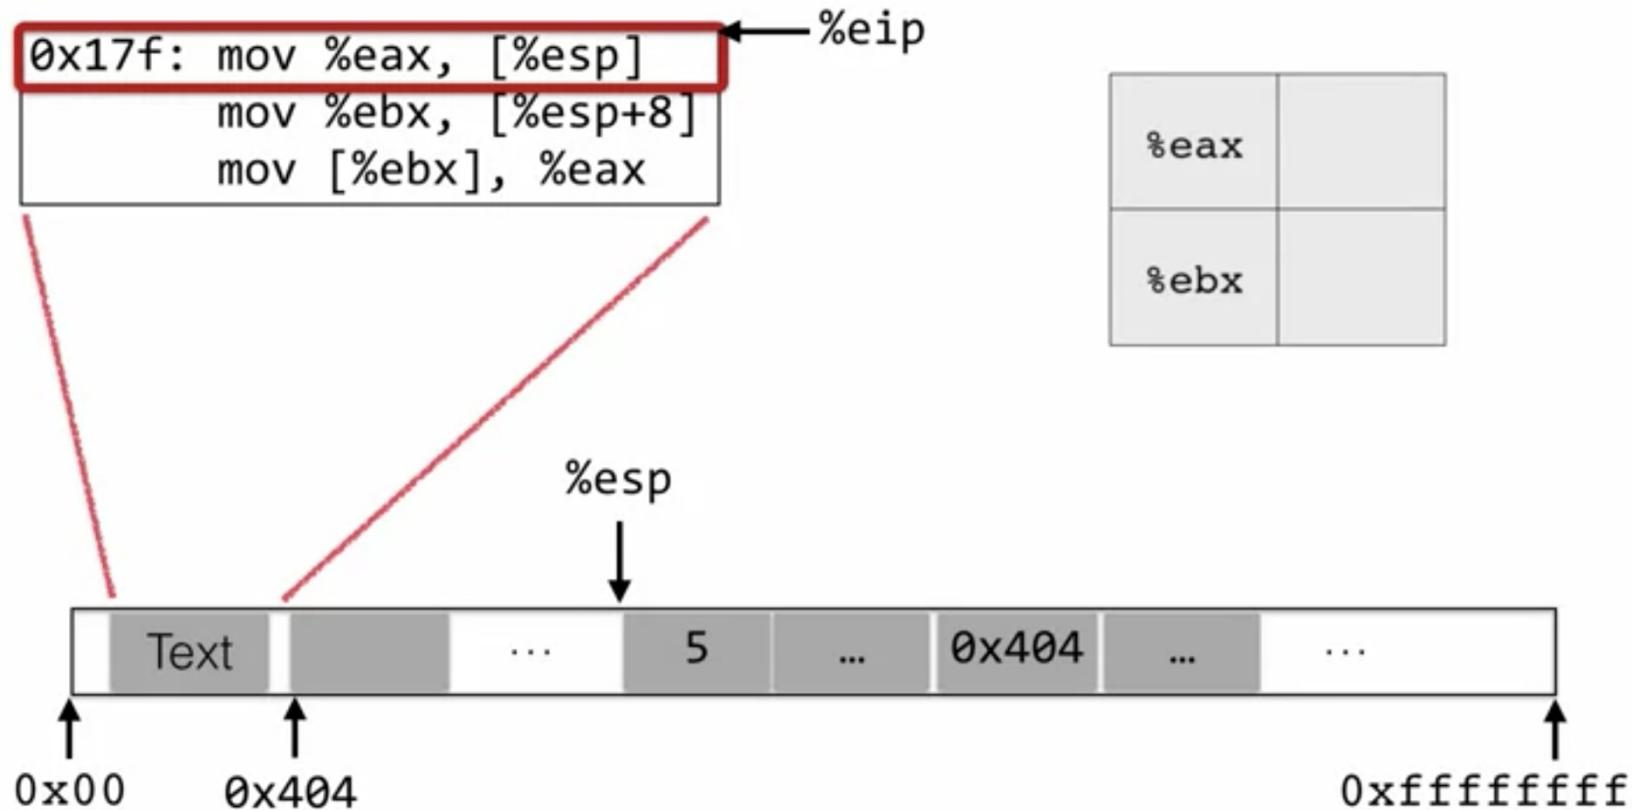
Simple ROP Example



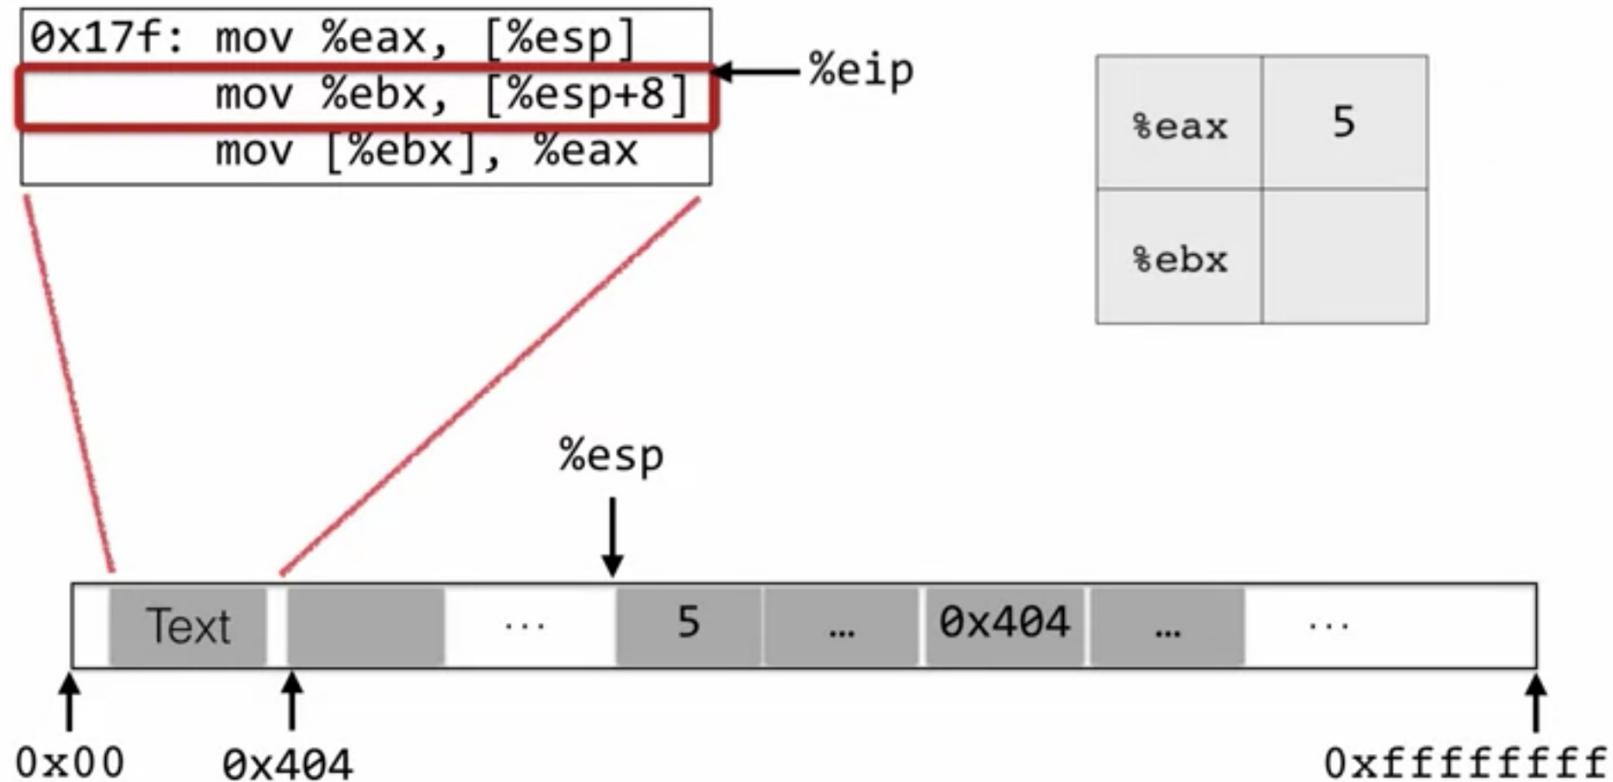
Simple ROP Example



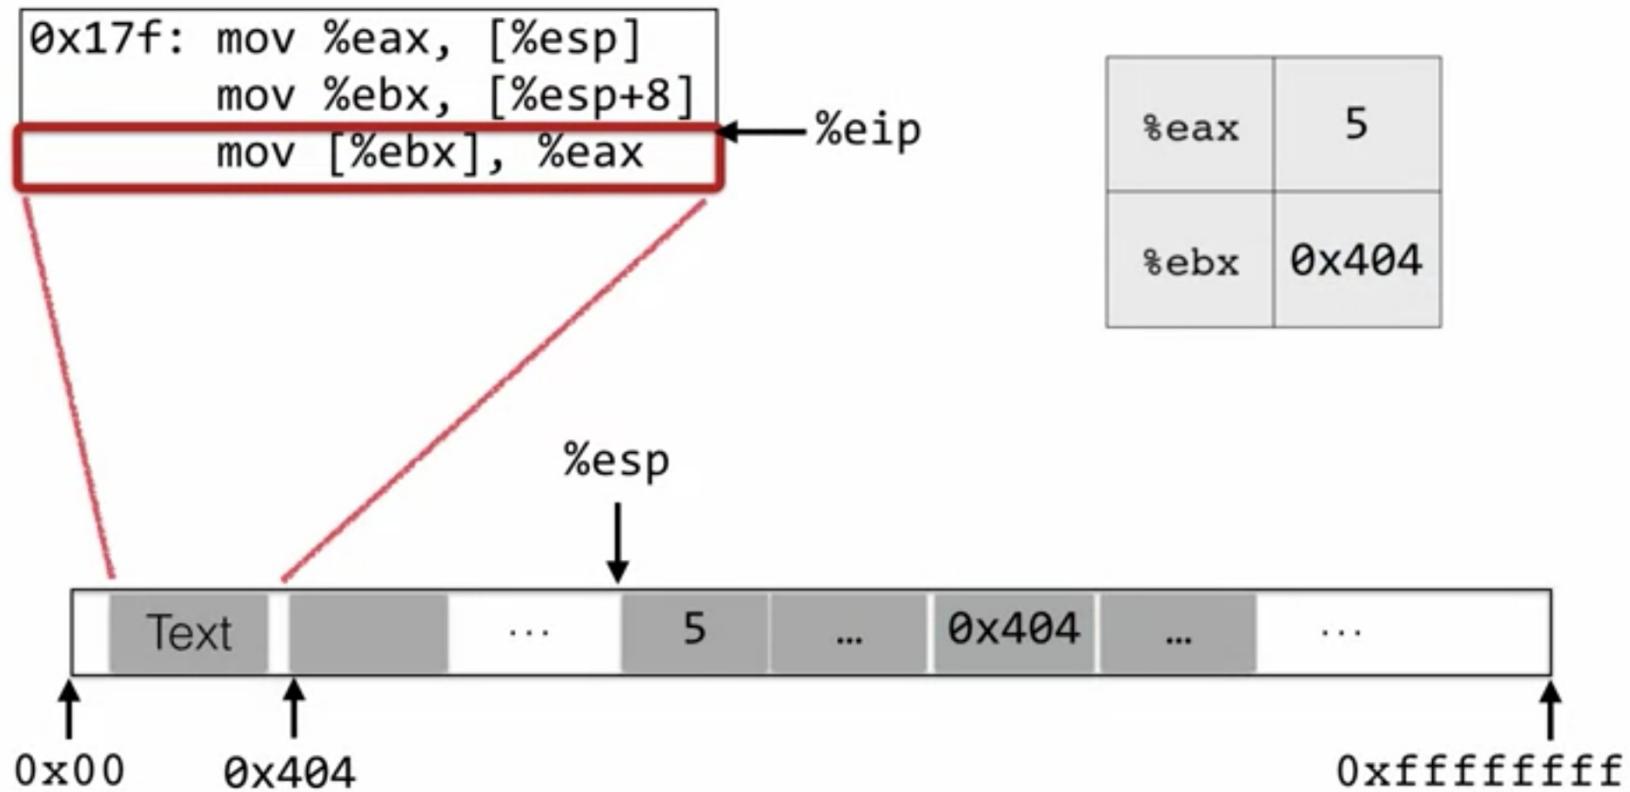
Another Example - Code Sequence



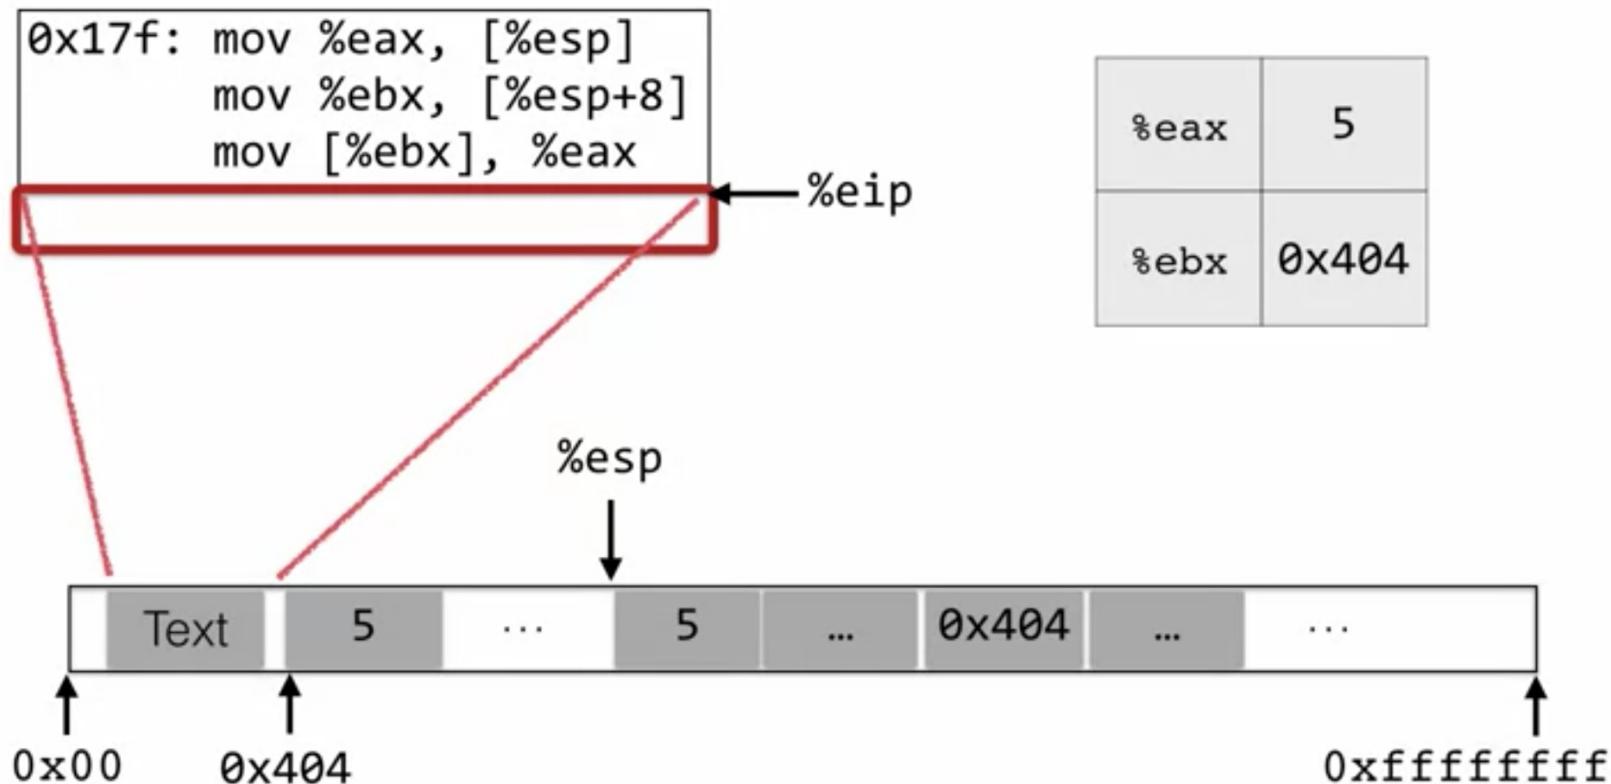
Another Example - Code Sequence



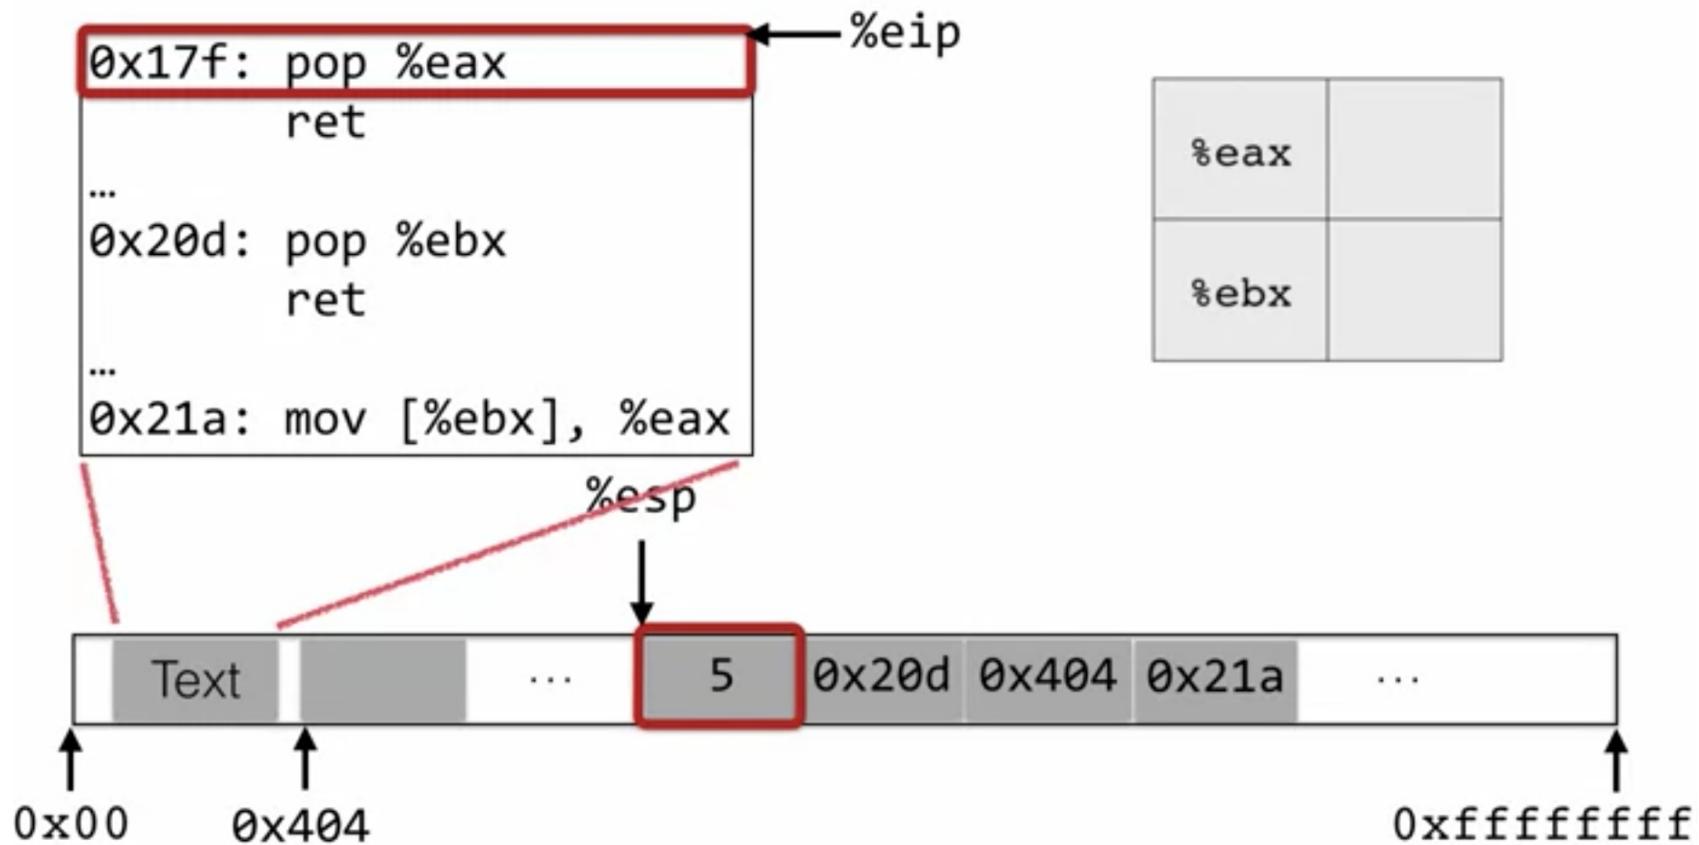
Another Example - Code Sequence



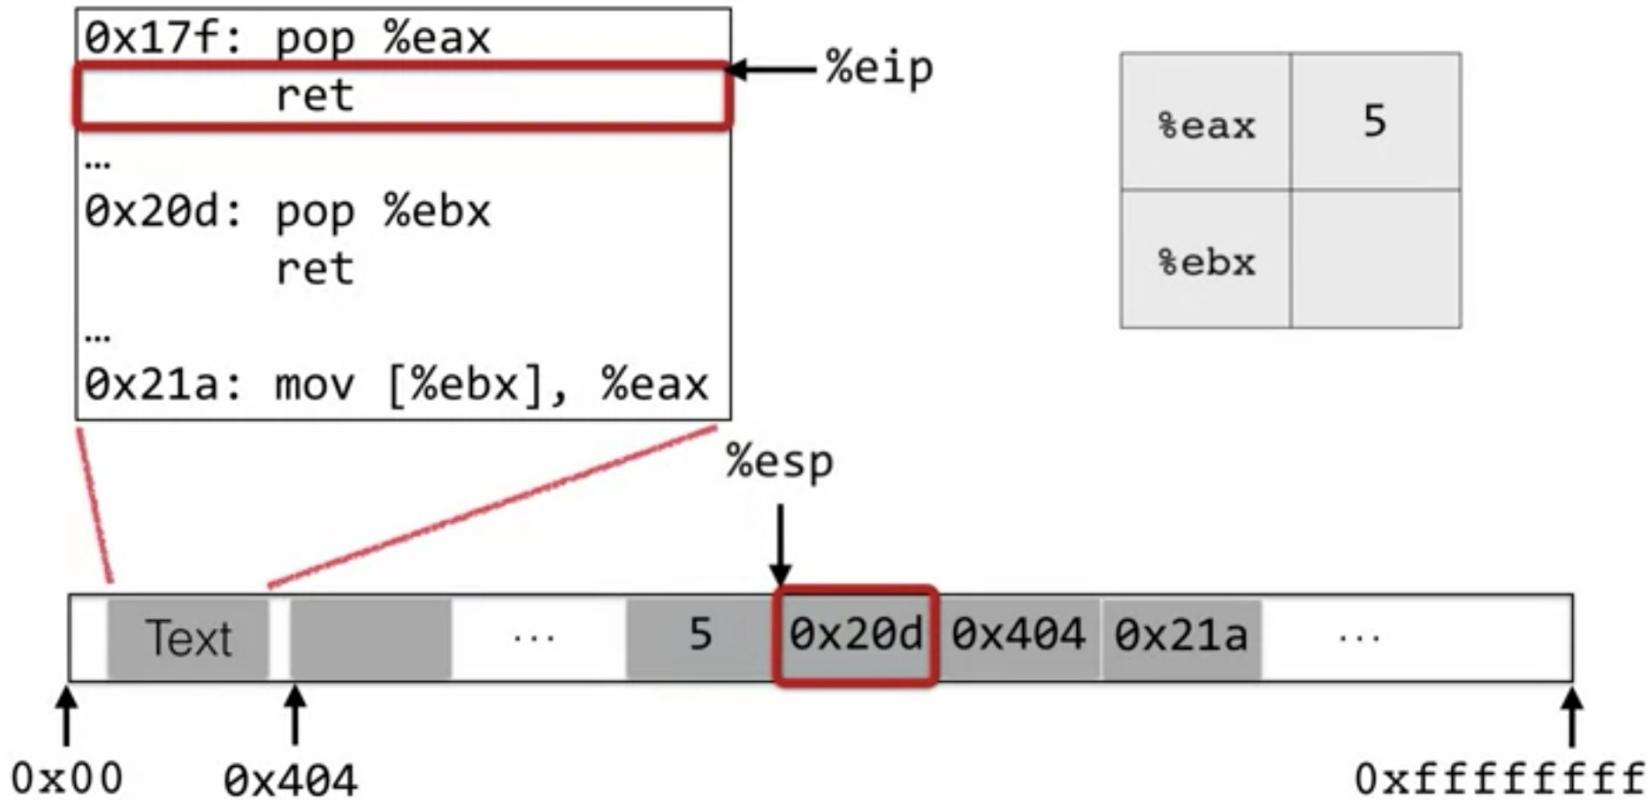
Another Example - Code Sequence



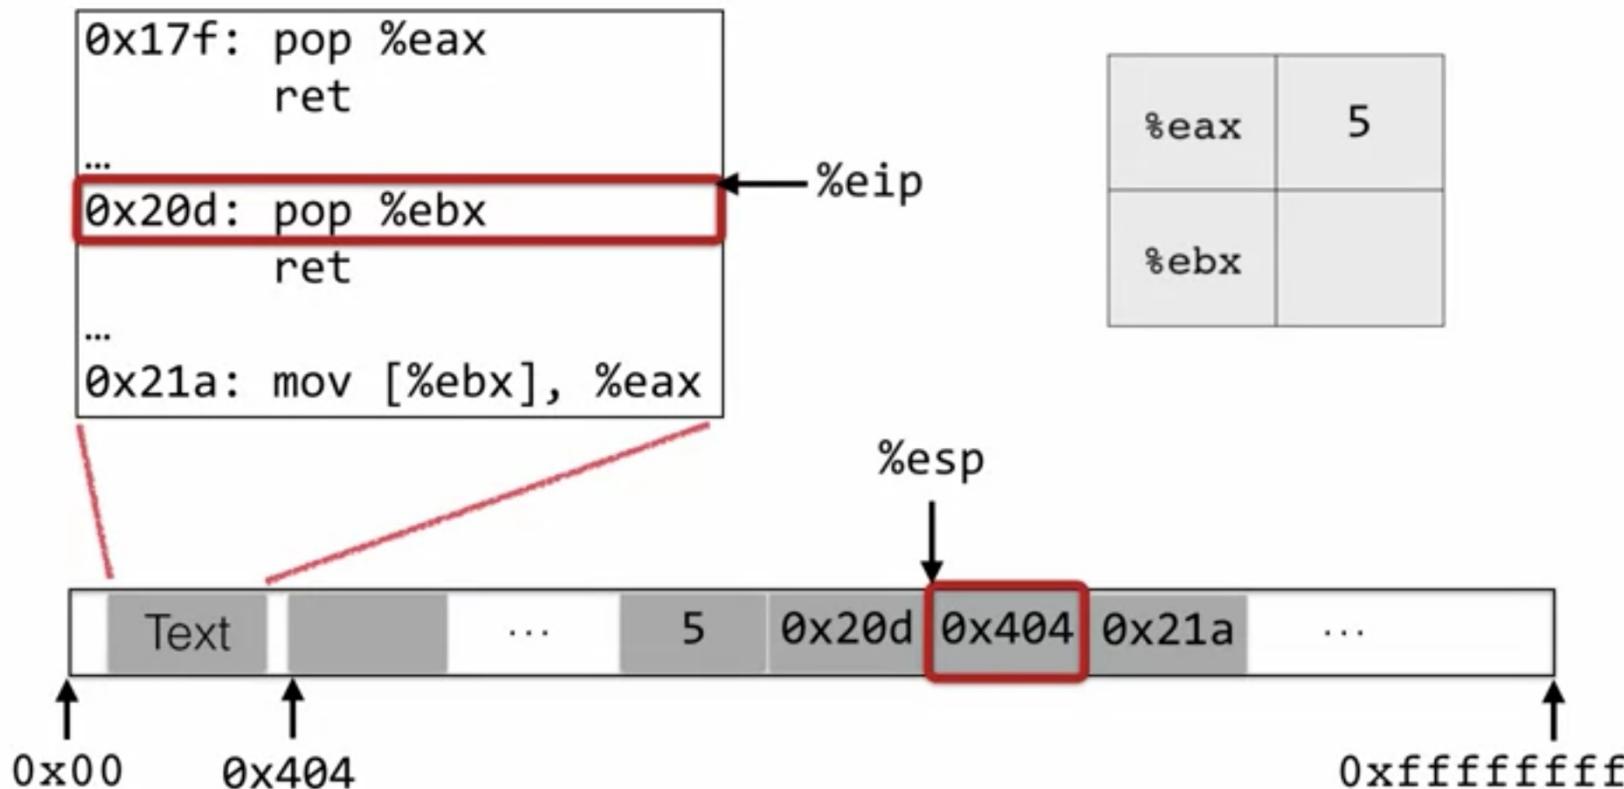
Another Example - ROP ~~Equivalence~~



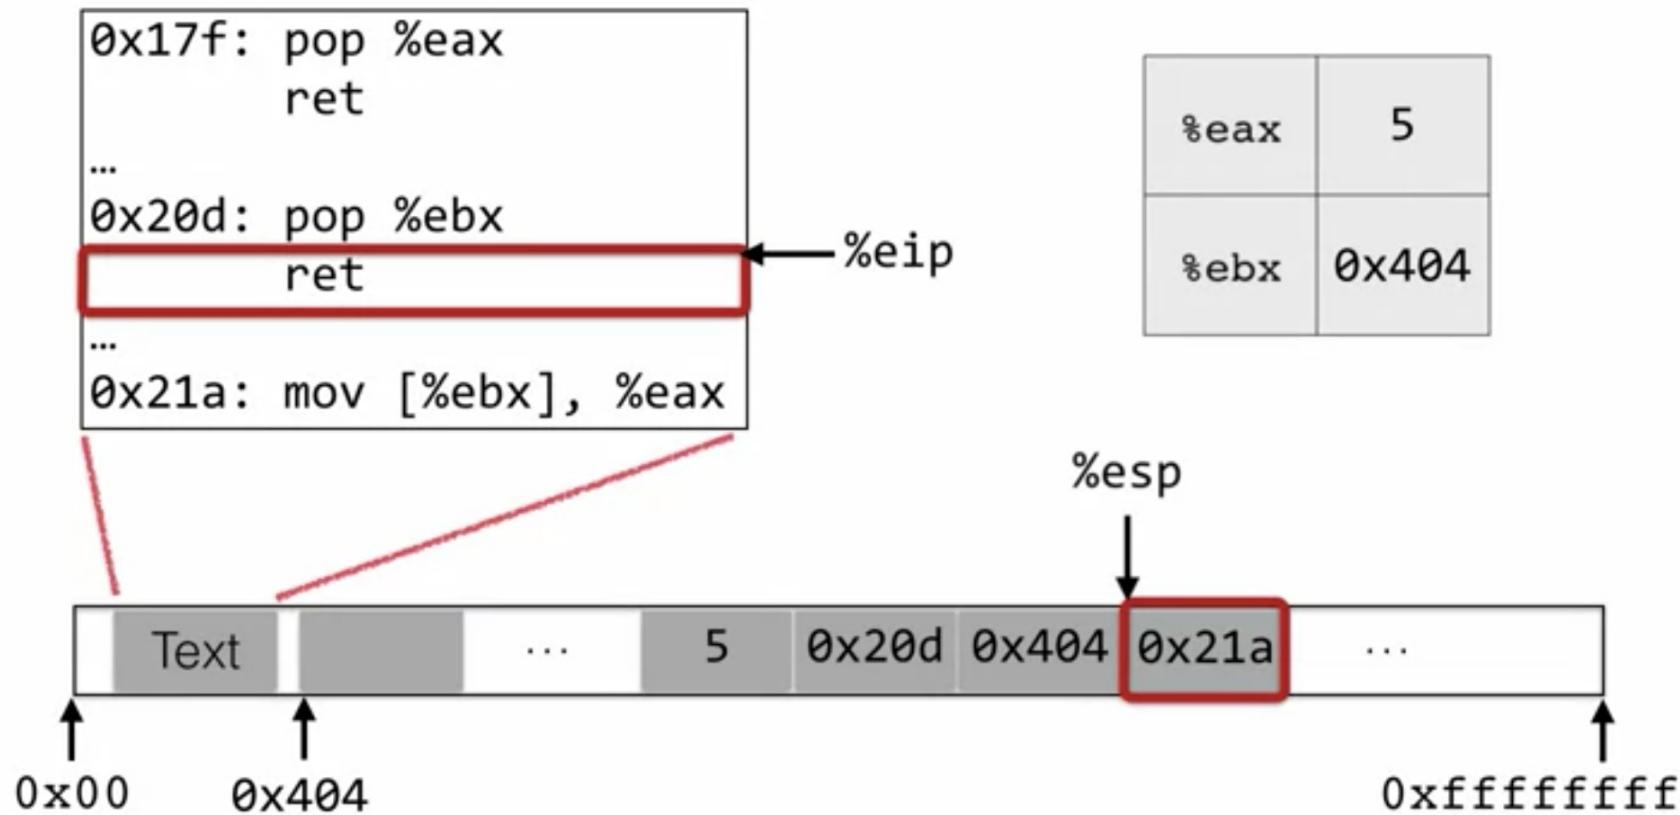
Another Example - ROP ~~Equivalence~~



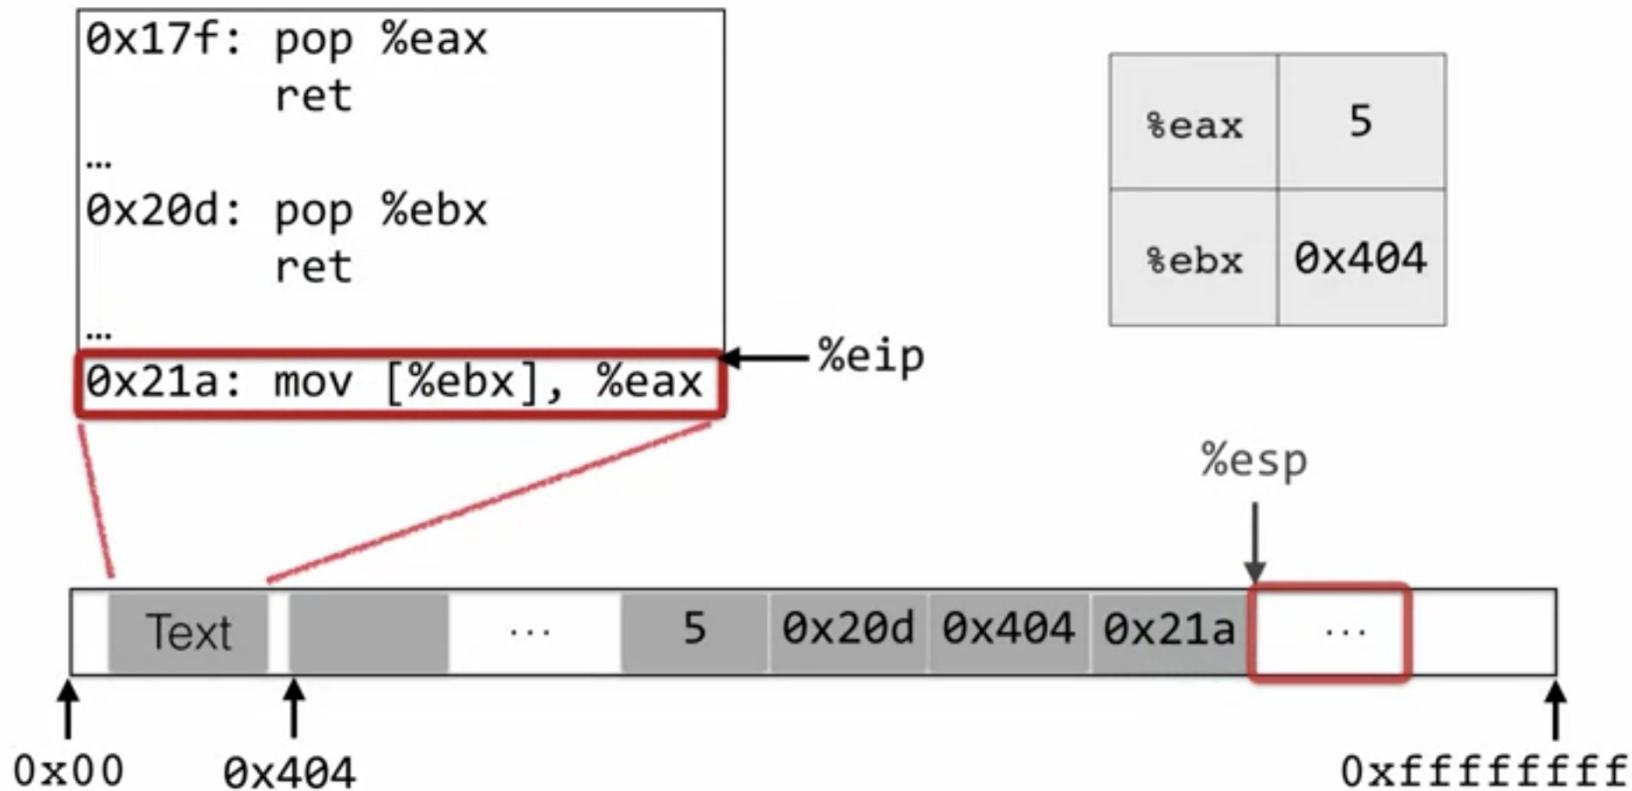
Another Example - ROP ~~Equivalence~~



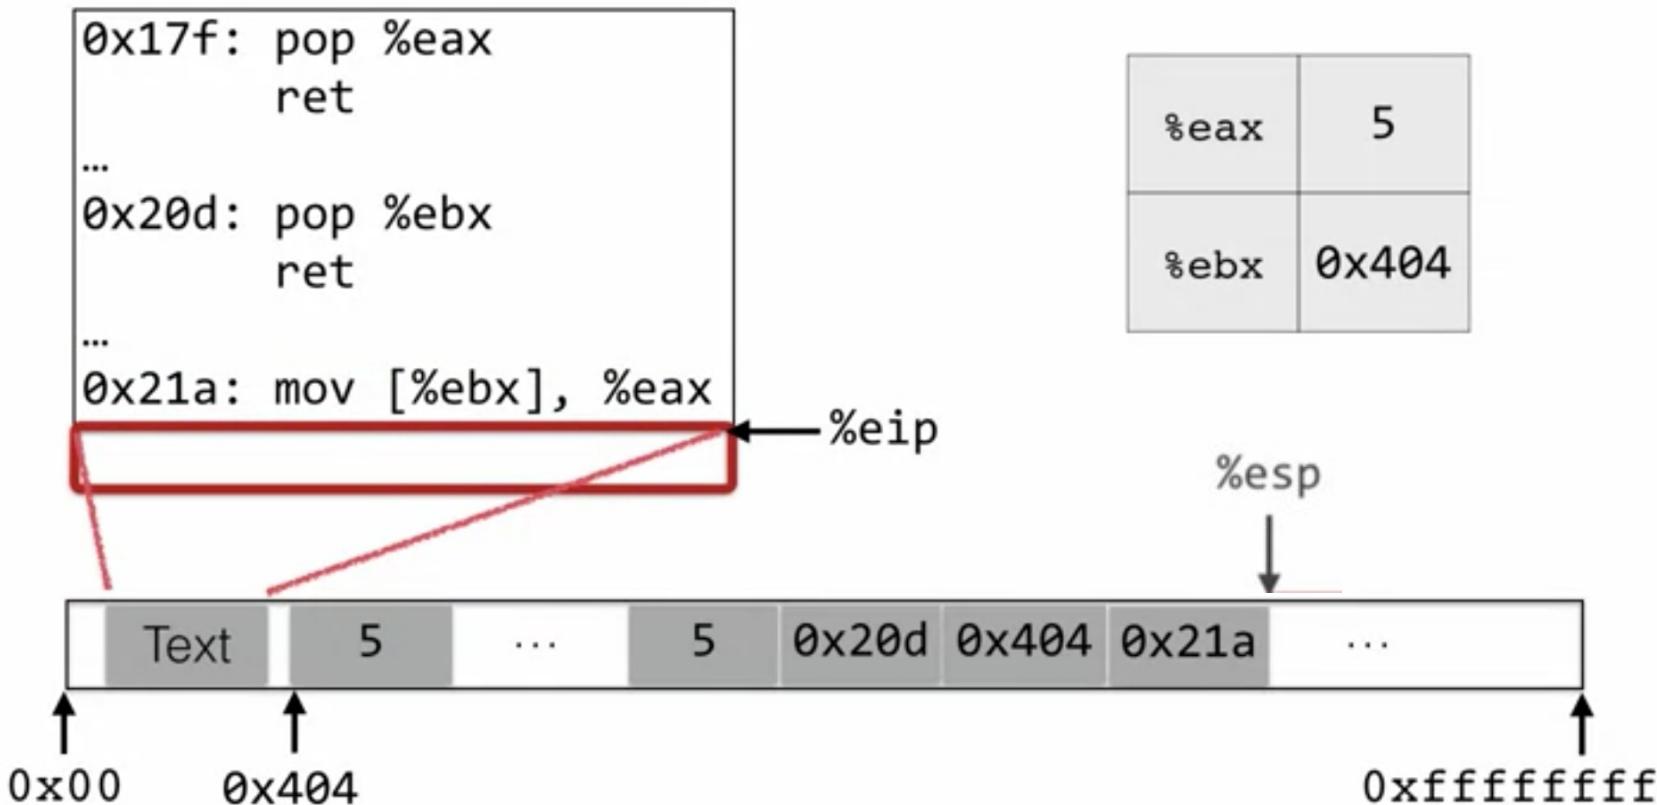
Another Example - ROP Equivalence



Another Example - ROP Equivalence



Another Example - ROP ~~Equivalence~~



Gadgets and How to Find Them

- ❖ Gadgets perform well-defined operations, e.g., load, xor, jump
- ❖ Each gadget makes use of one or more sequences of instructions found in text segment and ends with ret
- ❖ To find the sequences: first find a ret and store it in the root of a trie, then scan backwards to find children
- ❖ The “child-of” relation in the trie means that the child instruction immediately precedes the parent instruction in text segment
- ❖ Example: If in a trie, a node representing pop %eax is a child of the root node (ret), then we have discovered the sequence pop %eax; ret

Limitation of ROP

- ❖ Bypasses W^X
- ❖ But not stack canaries or ASLR
- ❖ Remains a real remote code injection threat

More Advanced ROP

- ❖ Blind ROP
- ❖ PIC ROP

Control Flow Integrity (CFI)

- ❖ Behavior-based detection
[Abadi et al., “Control-Flow Integrity,” (CCS’05)]
- ❖ Recall: stack canaries, W^X, and ASLR make one of the steps in stack smashing harder, but may not fully stop the attack
- ❖ Idea of CFI: observe the program’s behavior – Is it as expected?
- ❖ Challenges:
 - Define “expected behavior”
 - Detect deviations from the behavior efficiently
 - Avoid compromise of detector

Control Flow Integrity (CFI)

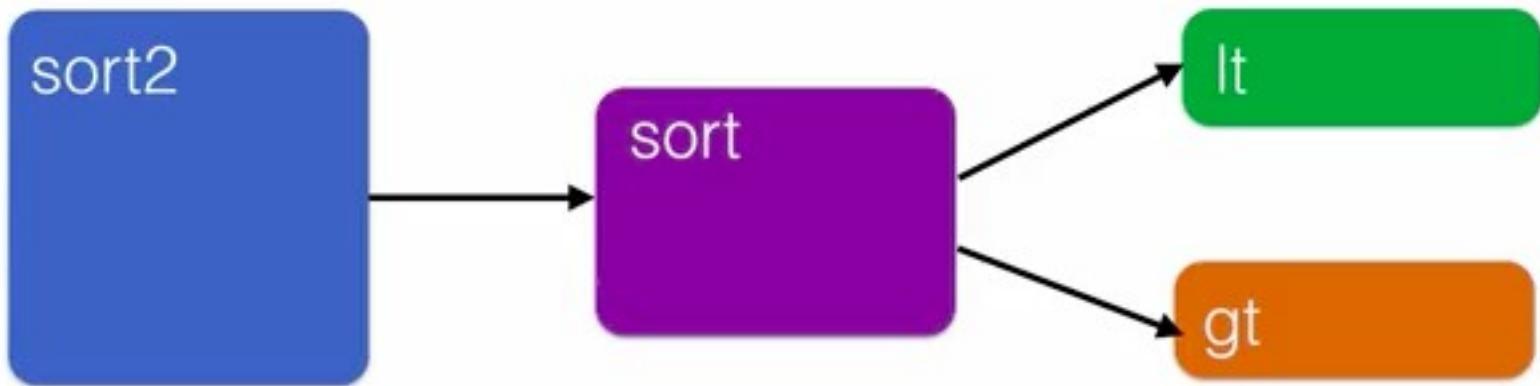
- ❖ Define “expected behavior”
 - control flow graph (CFG)
- ❖ Detect deviations from the behavior efficiently
 - Inline reference monitor (IRM)
- ❖ Avoid compromise of the detector
 - Sufficient randomness of labels; immutable code; non-executable data

Control Flow Graph (CFG)

- ❖ Call graph: which function calls which other

```
sort2(int a[], int b[], int len)
{
    sort(a, len, lt);
    sort(b, len, gt);
}
```

```
bool lt(int x, int y) {
    return x < y;
}
bool gt(int x, int y) {
    return x > y;
}
```



Control Flow Graph (CFG)

- ❖ CFG: breaks into basic blocks of machine code instructions
- ❖ Distinguish call from return
- ❖ Each block ends with some control transfer such as:

```
sort2(int a[], int b[], int len)
{
    sort(a, len, lt);
    sort(a, len, gt);
}
```

```
bool lt(int x, int y) {
    return x<y;
}
bool gt(int x, int y) {
    return x>y;
}
```

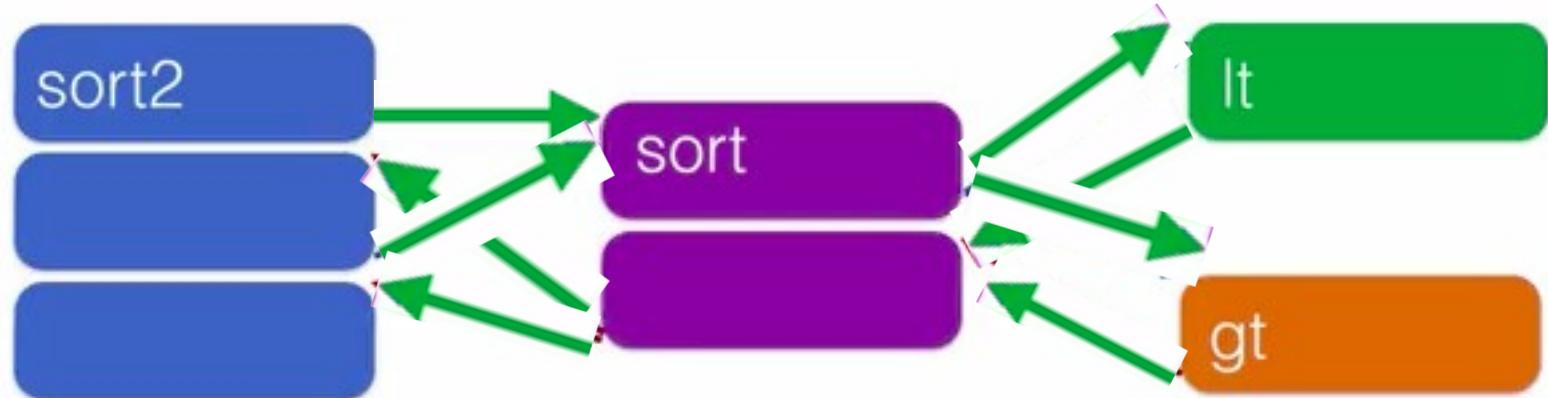


Control Flow Graph (CFG)

- ❖ Computed the call/return CFG in advanced
- ❖ Using program analysis (source code or binary)

```
sort2(int a[], int b[], int len)
{
    sort(a, len, lt);
    sort(a, len, gt);
}
```

```
bool lt(int x, int y) {
    return x<y;
}
bool gt(int x, int y) {
    return x>y;
}
```



Inline Reference Monitor (IRM)

- ❖ Checks for compliance with CFG
- ❖ No need to monitor direct calls: target address cannot be changed assuming code is immutable

```
sort2(int a[], int b[], int len)
{
    sort(a, len, lt);
    sort(a, len, gt);
}
```

```
bool lt(int x, int y) {
    return x<y;
}
bool gt(int x, int y) {
    return x>y;
}
```



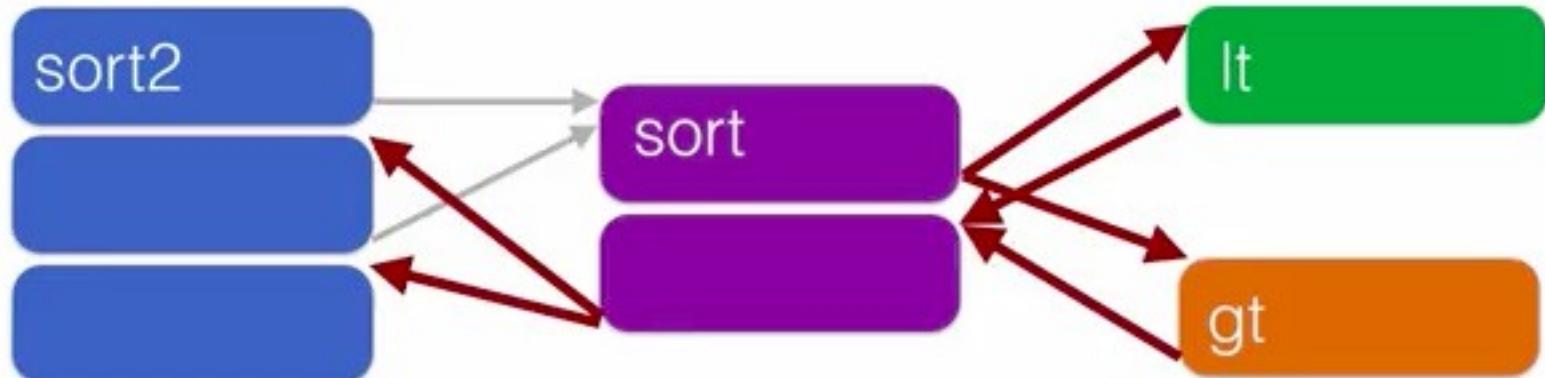
Direct calls (always the same target)

Inline Reference Monitor (IRM)

- ❖ Only monitors indirect transfer to improve efficiency: jmp, call via register, return

```
sort2(int a[], int b[], int len)
{
    sort(a, len, lt);
    sort(a, len, gt);
}
```

```
bool lt(int x, int y) {
    return x<y;
}
bool gt(int x, int y) {
    return x>y;
}
```

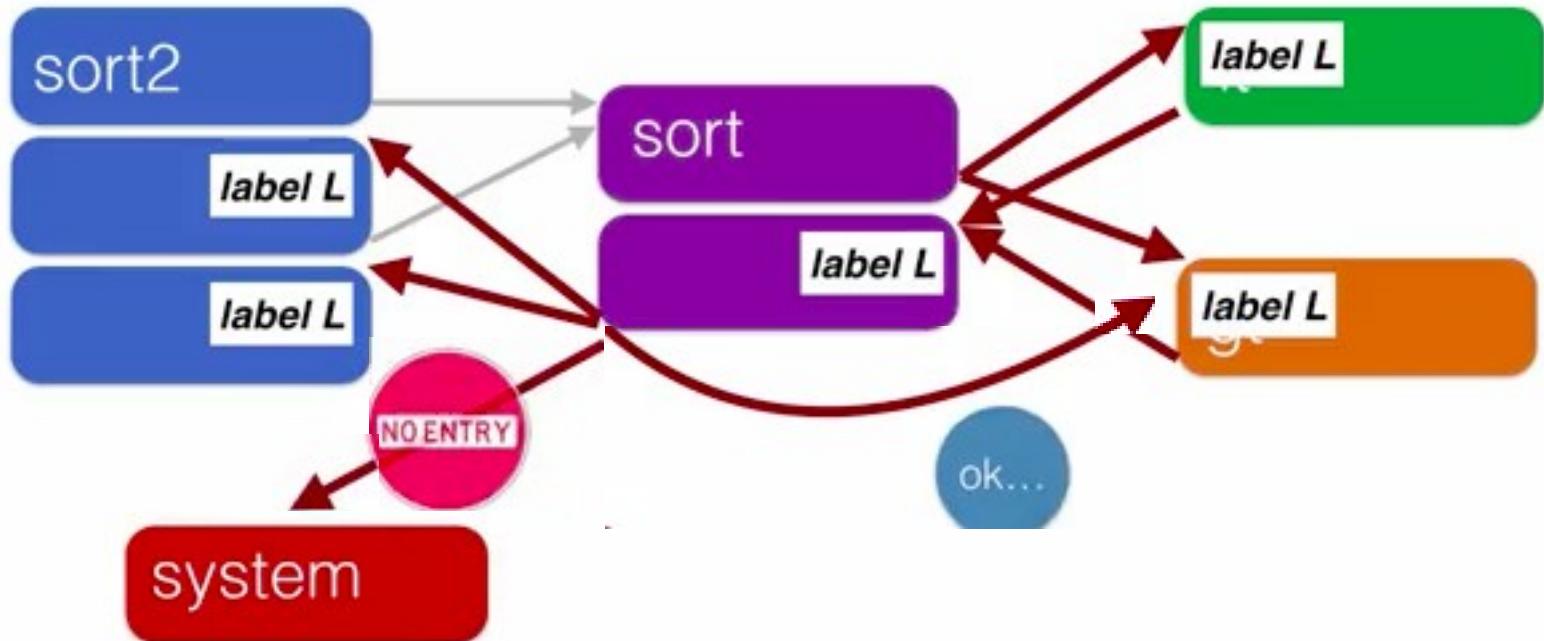


Indirect transfer (call via register, or ret)

Inline Reference Monitor (IRM)

- ❖ Implement IRM inline via machine code rewriting
- ❖ Insert a **label** just before the target address of an indirect transfer
- ❖ Insert **code** to check the label of the target at the source of the indirect transfer

Simple Labeling

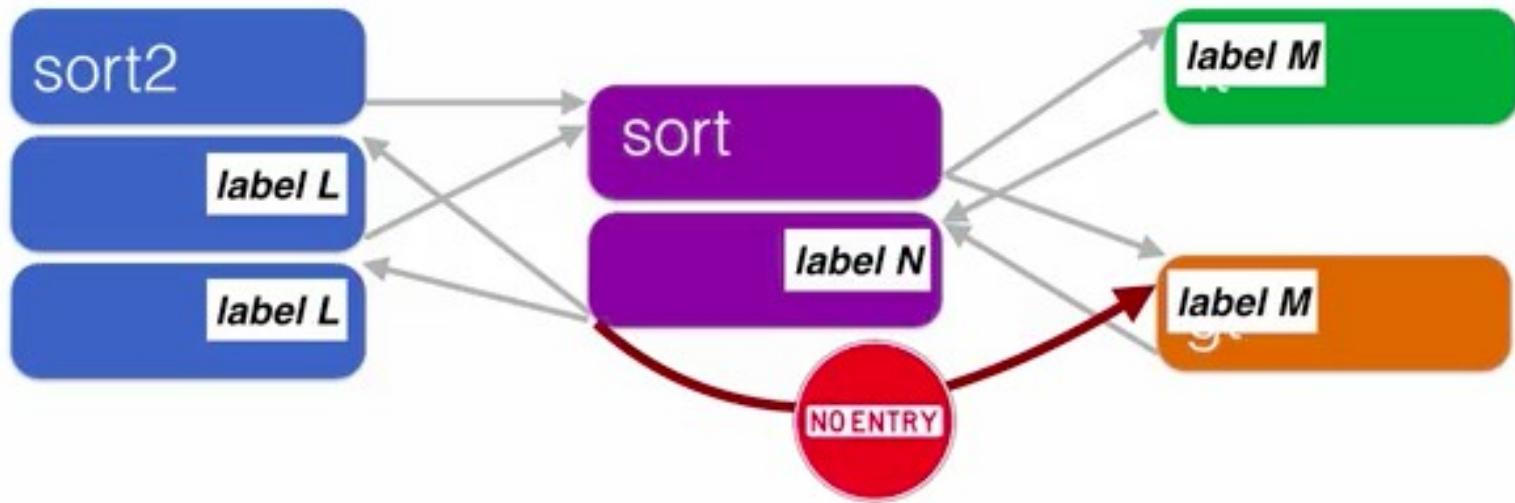


Use the same label at all targets

Blocks return to the start of direct-only call targets

but not incorrect ones

Detailed Labeling



Constraints:

- return sites from calls to **sort** must share a label (*L*)
- call targets **gt** and **lt** must share a label (*M*)
- remaining label unconstrained (*N*)

CFI Instrumentation

Opcode bytes	Source			Destination		
			Instructions			Instructions
FF E1		jmp	ecx		; computed jump	
				8B 44 24 04	mov	eax, [esp+4] ; dst
				...		
	can be instrumented as (a):					
81 39 78 56 34 12	cmp	[ecx]	, 12345678h	78 56 34 12	; data 12345678h	; ID
75 13	jne	error_label		8B 44 24 04	mov	eax, [esp+4] ; dst
8D 49 04	lea	ecx	, [ecx+4]	...		
FF E1	jmp	ecx				

Can We Defeat CFI?

- ❖ Inject code that has a legal label
 - won't work because we assume non-executable data
- ❖ Modify code labels to allow desired control flow
 - won't work because we assume code is immutable
- ❖ Modify stack during a check to make it seem to pass
 - won't work because we assume attacker cannot change the registers where we load relevant data (no TOCTOU)

Performance of CFI

- ❖ Overhead due to insertion of labeling and instructions that check for CFI compliance
 - classic CFI: 16% overhead average, 45% worst case
 - modular CFI: 5% average, 12% worst case
- ❖ MCFI can eliminate 95.75% of ROP gadgets by ruling their use non-compliant with CFG
- ❖ CFI achieves >99% average indirect-target reduction (AIR): percentage of possible targets of indirect jumps CFI rules out

Limitations of CFI

- ❖ CFI defeats control-flow modifying attacks
 - remote code injection; return-to-libc/ROP, etc.
- ❖ Does not prevent manipulation of flow allowed by CFG
- ❖ Does not prevent data leaks or corruptions (no control flow hijacking)
 - Heartbleed bug
 - “authenticated” flag
- ❖ But it is very effective against remote code injection and ROP

```
void func(char *arg1)
{
    int authenticated = 0;
    char buffer[4];
    strcpy(buffer, str);
    if(authenticated) { ...
}
```

Recap: Cat and Mouse

Attack: stack smashing

Defense: stack canaries (e.g., StackGuard), W^X

Attack: guess or bypass stack canaries, return-to-libc/ROP

Defense: ASLR (against return-to-libc/ROP)

ROP

Attack: brute force or information leak on ASLR, blind

Defense: CFI

CFI not mature, not effective against non-code injection

...

Defense Aategory 2: Avoid The Flaw Entirely

- ❖ Secure coding in C
 - Since the language provides few guarantees, developers must use discipline
 - In general, follow rules (good security practices) for good C coding
 - Good reference guide: CERT C coding standard
- ❖ Combined with advanced code review and testing
 - e.g., program analysis, fuzzing
- ❖ Use memory-safe languages
 - e.g., avoid C and C++, make C and C++ memory safe

Rule: Enforce Input Compliance

```
int main() {
    char buf[100], *p;
    int i, len;
    while (1) {
        p = fgets(buf, sizeof(buf), stdin);
        if (p == NULL) return 0;
        len = atoi(p);
        p = fgets(buf, sizeof(buf), stdin);
        if (p == NULL) return 0;
        for (i=0; i<len; i++)
            if (!iscntrl(buf[i])) putchar(buf[i]);
            else putchar('.');
        printf("\n");
    }
}
```

Recall earlier example

- } Read integer
- } Read message
- } Echo back (partial) message

Rule: Enforce Input Compliance

```
int main() {
    char buf[100], *p;
    int i, len;
    while (1) {
        p = fgets(buf, sizeof(buf), stdin);
        if (p == NULL) return 0;
        len = atoi(p);
        p = fgets(buf, sizeof(buf), stdin);
        if (p == NULL) return 0;
        for (i=0; i<len; i++)
            if (!iscntrl(buf[i])) putchar(buf[i]);
            else putchar('.');
        printf("\n");
    }
}
```

*May exceed
actual message
length!*

*Recall earlier
example*

} Read integer
} Read message
} Echo back
 (partial)
 message

Rule: Enforce Input Compliance

```
int main() {
    char buf[100], *p;
    int i, len;
    while (1) {
        p = fgets(buf, sizeof(buf), stdin);
        if (p == NULL) return 0;
        len = atoi(p);
        p = fgets(buf, sizeof(buf), stdin);
        if (p == NULL) return 0;
len = MIN(len,strlen(buf));
        for (i=0; i<len; i++)
            if (!iscntrl(buf[i])) putchar(buf[i]);
            else putchar('.');
        printf("\n");
    }
}
```

Recall earlier example

Sanitizes input
to be compliant

Rule: Enforce Input Compliance

```
char digit_to_char(int i) {  
    char convert[] = "0123456789";  
    return convert[i];  
}
```

Rule: Enforce Input Compliance

```
char digit_to_char(int i) {
    char convert[] = "0123456789";
    if(i < 0 || i > 9)
        return '?';
    return convert[i];
}
```

- ❖ Unfounded trust in received input is a recurring source of vulnerability

Rule: Use Safe String Functions

```
char str[4];
char buf[10] = "fine";
strcpy(str,"hello"); // overflows str
strcat(buf,"day to you"); // overflows buf
```

```
char str[4];
char buf[10] = "fine";
strlcpy(str,"hello",sizeof(str)); //fails
strlcat(buf,"day to you",sizeof(buf)); //fails
```

Replacements

- ... for string-oriented functions
 - `strcat` ⇒ `strlcat`
 - `strcpy` ⇒ `strlcpy`
 - `strncat` ⇒ `strlcat`
 - `strncpy` ⇒ `strlcpy`
 - `sprintf` ⇒ `snprintf`
 - `vprintf` ⇒ `vsnprintf`
 - `gets` ⇒ `fgets`
- Microsoft versions different
 - `strcpy_s`, `strcat_s`, ...

Rule: Don't Forget NUL Terminator

```
char str[3];
strcpy(str,"bye"); // write overflow
int x = strlen(str); // read overflow
```

```
char str[3];
strlcpy(str,"bye",3); // blocked
int x = strlen(str); // returns 2
```

Rule: Understand Pointer Arithmetic

```
int buf[SIZE] = { ... };
int *buf_ptr = buf;

while (!done() && buf_ptr < (buf + sizeof(buf))) {
    *buf_ptr++ = getnext(); // will overflow
}
```

```
while (!done() && buf_ptr < (buf + SIZE)) {
    *buf_ptr++ = getnext(); // stays in bounds
}
```

(Better) Rule: Favor Safe Libraries

- ❖ Libraries encapsulates well-thought-out design
- ❖ Networking: Google protocol buffers, Apache Thrift
 - for dealing with network transmitted data
 - ensure input validation, parsing, etc.
 - efficient

Defense Summary

- ❖ Defense category 1 and 2 are complementary

Other Programming Oversights

- ❖ TOCTOU
- ❖ Off-by-one error
- ❖ Undocumented entry points in programs
- ❖ ...
- ❖ Not all programming oversights have security implications

TOCTOU (Race Condition)

- ❖ Time-of-Check to Time-of-Use
 - Between access check and use (race window), data can be altered

```
function withdraw($amount)
{
    $balance = getBalance();
    if ($amount <= $balance) {
        $balance = $balance - $amount;
        echo "You have withdrawn: $amount";
        saveBalance($balance);
    }
    else {
        echo "Insufficient funds.";
    }
}
```

TOCTOU (Race Condition)

- ❖ Time-of-Check to Time-of-Use
 - Between access check and use (race window), data can be altered

```
void silly_function(char *pathname) {
    struct stat f, we;
    int rfd, wfd;
    char *buf;
    stat(pathname, &f);
    stat("/what/ever", &we);
    if (f.st_dev == we.st_dev && f.st_ino == we.st_ino) {
        return;
    }
    rfd = open(pathname, O_RDONLY);
    buf = malloc(f.st_size - 1);
    read(rfd, buf, f.st_size - 1);
    close(rfd);

    stat(pathname, &f);
    if (f.st_dev == we.st_dev && f.st_ino == we.st_ino) {
        return;
    }
    wfd = open(pathname, O_WRONLY | O_TRUNC);
    write(wfd, buf, f.st_size - 1);
    close(wfd);
    free(buf);
}
```